

A Guide to Designing Experiments to Test Statistical Graphics

Emily Robinson*

Statistics Department

California Polytechnic State University

Heike Hofmann†

Statistics Department

Iowa State University

Susan Vanderplas‡

Statistics Department

University of Nebraska, Lincoln

Abstract

In this paper, we discuss considerations and methods for experimentally testing visualizations. We discuss levels of user engagement with graphics, common issues when developing a sampling or data generation model, the importance of pilot testing, and data analysis methods. Along the way, we also share hard-learned lessons on what not to do and some of the unique pitfalls of human testing in statistical and visualization research.

Index Terms: Human-centered computing—Visualization—Visualization design and evaluation methods ; Human-centered computing—Visualization—Empirical studies in visualization

Introduction

Charts and data visualizations are primary methods of communicating scientific information to the public in what we typically hope is an easy-to-digest, visually attractive package. There are many strategies for creating charts and graphs, from Tufte-esque minimalism [?] to charts designed with extra imagery and aesthetic appeal that draw the viewer’s attention and persist in memory [?, Ch 3]. For a specified type of data, there are also usually many different chart forms to display that data: for instance, if we have a set of categorical data and we wish to show the relative proportions of each category, we could do so using a stacked bar chart or the polar equivalent, a pie chart. There have been several attempts to list out all of the types of charts [?], create a taxonomy of charts [?, ?], and even to classify charts using a domain-specific grammar of graphics [?]. With all of the different design choices available, how are chart creators to know what is the best approach for communicating data to the appropriate audience?

There are heuristics, rules-of-thumb, and various guidelines [?, ?, ?, ?, ?] for creating useful and visually attractive charts, but as scientists, we believe empirical evidence is essential for making correct decisions. The best way to establish the efficacy of various design decisions, then, is to test our charts by having them evaluated by people under controlled conditions and gathering appropriate data [?, ?]. It can be extremely challenging to design studies which strike the right combination of internal and external validity. Simply asking people to read quantities off of a graph may not generalize beyond the questions asked or the data used in the chart [?, ?, ?, ?], but designing a study that is sufficiently robust to those issues requires manipulation of so many factors as to be functionally close to impossible to complete. In this paper, we attempt to distill lessons learned from several different sets of experimental graphics experiments [?, ?, ?, ?, ?], discussing

the use of several different testing methods for statistical graphics [?], the process of designing a graphical experiment, and analysis of the resulting empirical data. It is our hope that by writing this paper, we will prevent other researchers from repeating hard lessons in what to do (and what not to do), lowering the barriers that exist for conducting research involving empirical tests of statistical graphics.

Testing Methods and User Engagement

We will start with an overview of different testing methods that we have used to empirically assess graphics. While it is certainly possible to use equipment such as eye-tracking devices, here we will primarily focus on studies which can be conducted over the internet, because these experiments have the lowest overhead, offer relatively fast data collection, and provide useful results that can be very clear if the experiment is designed correctly. Our toolkit is R-based [?], and includes ggplot2 [?] and Shiny [?] as primary components, as well as JavaScript (including D3 [?]) to customize the Shiny interface for many of the experiments.

One primary concern when testing graphics is the level of user engagement required. For instance, if we are testing whether someone can detect an effect such as a trend line in noisy data, this is a *perceptual* question best addressed using methods which allow the user to interact with the data on a basic visual level. Sometimes, it may be possible to set up a scenario where the user adjusts the plot directly, as in our experiment testing the strength of the sine illusion experiment [?]: users adjusted the strength of the illusion correction until the lines appeared to be the same length, providing a direct measure of the magnitude of the sine illusion’s effect, as in Figure 1.

Graphical Cognition

Please adjust the value below until you feel the lines are (approximately) the same size. No visual aid or ruler should be necessary.

Adjustment Reset

When you are satisfied, press the "Finished" button for another question. Each question should take less than a minute to complete.

Finished Skip Question

Questions Answered: 0
Questions Skipped: 0 of 2 attempted

Amazon Worker ID

Only necessary for payment through Amazon Turk.

This experiment is IRB exempt. (More Information)




Figure 1: Direct adjustment of a plot in a perceptual task

Statistical lineups [?, ?] are a more generally applicable testing tool for perceptual questions such as “which chart displays this data more clearly” [?] while simultaneously assessing the significance of the graphical finding in the chart. As in the criminal procedure, lineups embed a target plot with signal into an array of 19 other innocent “null” plots that show only noise; if viewers consistently pick the target plot at a higher rate than the 19 null plots, then the target plot is said to be visually significant [?, ?] and a “see” value (the visual analogue of a p -value for a statistical test) can be

*email: erobin17@calpoly.edu

†email: hofmann@iastate.edu

‡email: svanderplas2@unl.edu

calculated [?]. In another variation of the statistical lineup procedure, two models are compared, with target plots from each model embedded in the array of 20 total plots; the remaining null plots are constructed from a mixture model that blends the two competing models. This variation allows the experimenter to assess graphical design choices to determine whether they effectively emphasize structural differences in the data and assign a hierarchy of visual attributes [?]. Statistical lineups do not require the experimenter to provide any contextual information for the data: all of the necessary information is embedded in the choice of a null statistical model, which is convenient for testing purposes, but does not require engagement at a level that would allow for testing the viewer’s understanding of the information shown in the chart. An example of a statistical lineup for testing the perception of exponential growth is shown in Figure 2.

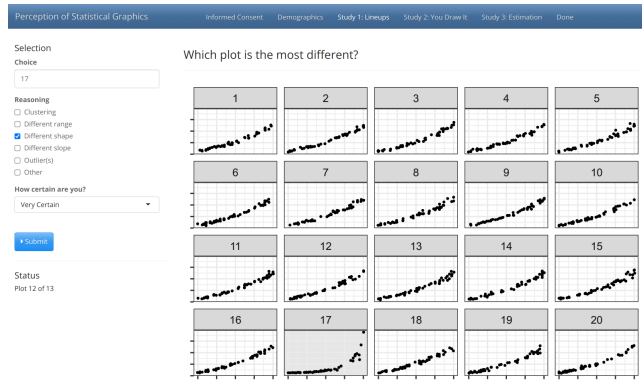


Figure 2: A lineup experiment testing whether users can perceive differences in rates of exponential growth on linear and log scales. In addition to asking users to select the plot which is the most different, we also asked for basic reasoning and the user’s level of confidence in their answer.

To assess the viewer’s *understanding* of information shown in a chart, we must ask questions and allow the user to provide feedback. User feedback may be collected on a numerical scale or through the use of written comments, recorded “think-aloud” processes, and other more qualitative interaction methods. We have conducted studies asking users to provide us with a written description of what the chart shows within a wider story-driven scenario (Figure 3), and we have also asked for written feedback about why a certain panel of a statistical lineup was chosen; the insights from participant responses to these questions have been rich and insightful (and will be discussed more in the analysis section). While we have not (to date) recorded users talking out loud about what they are seeing, it is entirely feasible that so-called “think aloud” processes could be used within a Shiny application, with audio recordings saved to the server for transcription and analysis [?, ?, ?].

Of course, as we have confined ourselves to the discussion of online, asynchronous testing methods, every user interaction with the testing materials (typically hosted on a web page) can also be recorded along with time stamps, mouse positions, browser size and screen resolution, and other information. While we have not used this type of information heavily in our experimental analyses thus far, we do collect some time stamp data in order to assess how long participants spend on each question. What we typically see, is that the first round of test questions takes the longest for participants to




Figure 3: This question asks users to write out a description of how the population of Ewoks changes over time, without any further cues, to determine whether participants default to multiplicative or additive language descriptions.

complete. Additional experiments do not affect accuracy (i.e. there is no immediate learning effect) until after ‘too many’ tests cognitive fatigue proves to be detrimental to accuracy [?]. This sweet spot between user replicates and fatigue depends on the cognitive burden in each test and should factor into designing the experiment. In addition, we have also provided participants with supportive tools, such as “scratch pads” and limited function calculators built into the Shiny application that allow us to gain additional insight into the estimation and visual comparison process when more complex calculations are required to answer the numerical estimation question (Figure 4). These supportive tools have provided a wealth of information that, while not easy to analyze and generally messy, provides real insight into how participants were looking at the data, what strategies they tried and discarded for reading the chart, and what visual estimation methods were used.




Figure 4: This question asks participants for a numerical estimate, but provides a basic calculator and scratchpad. All user interactions with the calculator and scratchpad are logged, providing insight into the user’s thought process and estimation strategy.

One of the most difficult components of designing an experiment which asks users to directly estimate information from a chart using a full scenario (background information, etc. as well as contextual details from the chart) is that the questions must be extremely carefully constructed. Mathematics education researchers provide guidelines for selecting different levels of questioning in order to assess graph comprehension: literal reading of the data, reading between the data, and reading beyond the data [?, ?, ?, ?]. In a recent study, we identified questions based on this framework to

evaluate direct estimates and extend those estimates to make comparisons between two points.

Even when great care is taken with the construction of the question, participant answer accuracy is fundamentally limited by the fact that many participants do not read and interpret the question with the care and precision that it was written. Questions that ask participants to e.g. estimate the multiplicative change in a quantity at two time points may be misunderstood as asking for an estimate of the additive difference, and the resulting estimates are then one or more orders of magnitude off of the correct answer. This is one area where lineup tests are convenient - they do not depend on participants to understand the nuances of language or scenarios built around the chart under investigation. However, in some situations it may be sufficient to ask participants to estimate direct numerical quantities that have little contextual information, as we did in [?] when assessing the accuracy of framed plots re-created from the Statistical Atlas.

It is also possible to design a task that requires participants to *engage directly* with an interactive visualization. This can be useful in a directed task, where users are asked to interact with the chart in a specific way and the result is recorded, but it is also possible to use interactive visualizations in an open-ended task, recording how users engage with the graphic in an exploratory (as opposed to goal-directed) manner. In one recent experiment, we asked participants to forecast an exponential trend, with data presented on either a linear or log scale. Using JavaScript code modified from New York Times interactive graphics “You Draw It” features [?], we had users draw trend lines with their computer mouse and make forecasts directly on interactive charts, with the data and user-drawn predictions recorded to our database. With interactive graphics rendered using JavaScript (or other web libraries), the only limit to the types of questions one can ask in testing graphics is one’s ability to write code to interact with the visualization library. This type of testing method can be extremely natural for participants, but it also is hard to generalize when discussing testing methods because of the potential range of applications where it might be employed.

Whichever testing method is chosen should be appropriate to the type of question under investigation and the level of visual and cognitive engagement required to answer that question. While lineups are excellent tools for assessing perceptual questions, they cannot address questions aimed at understanding how people use charts within the wider context of a story or practical task; this requires more direct methods with higher ecological validity.

Lineups require a great amount of work in developing a null sampling model which replicates the important visual features of the data, but model specification is extremely important when asking users questions about any simulated data set or even when examining user interactions directly. It is important to carefully calibrate the model parameters to allow for some variability, but not too much, and to ensure that participants can succeed at the task and do not feel like they are being made to analyze random noise. This Goldilocks-style problem is the focus of the next section.

Developing a Model

Once you have decided what graphical task you want to explore or test, it is necessary to develop a model which can be used to test the graphical features you want to explore in a precise manner. This is the single longest part of the entire experimental design and execution process, in part because

choosing a model that replicates important visual features of the data is an extremely complex process [?, ?, ?].

There are two main options when developing a statistical model for graphical testing: you can either start with a large dataset and sample from that dataset [?], or you can start from a model and sample data from that model generating process [?, ?, ?]. This decision is largely determined by the availability of a large dataset containing the requisite features of interest and the qualities being manipulated in the experiment. For instance, Hofmann et al. [?] used samples of different sizes from a pre-existing dataset to manipulate the amount of signal in each comparison; with a small sample, there is less signal and the same amount of noise, making the true plot harder to spot. In many situations, though, a convenient dataset with the right properties is harder to acquire, and it becomes necessary to develop a sampling model to generate data for user evaluation.

The tools we discuss in the remainder of this section can be applied both to pre-existing data sets and to model-based sampling methods. We will do our best to keep this discussion general, but we will provide details in some situations to ensure that our general descriptions are not hopelessly vague.

Screening Parameters with Simulation

Visual assessment is key, but narrowing the parameter space down to something that can be statistically determined is useful for providing a sort of numerical summary estimate of potential visual difficulty. If using model based sampling methods, it is important to iteratively assess the sampling procedure through simulation: Start with a sampling method and define a parameter space that is appropriate for the model, and then simulate many data sets from each possible parameter combinations using a grid search. When sampling from a large data set, it is still important to assess the sample size and any potential stratification methods employed in order to ensure that the visual effect is represented.

In both cases, it can help to employ some general feature engineering - what numerical statistic provides an estimation of the basic concept which is being visually evaluated? For instance, we have used

- R^2 as a measure of the strength of a linear relationship
- Gini inequality as a measure of the strength of clustering
- lack-of-fit statistics to assess the amount of curvature in an exponential relationship (shown in Figure 5)

Then, a wide range of potential combinations of parameter values or sampling strategies can be explored and summarized graphically; if the numerical statistic cannot differentiate between the null and target under a condition, it is reasonable to expect that a visual inspection of the data may also not show significant results. As with any measure, it is important that difficulty levels span a range from easy to hard; we do not learn anything from finding out that everyone can distinguish all of the combinations.

Fine-Tuning Parameter Choices

Once an appropriate set of parameters are identified using the numerical screening method, then it is important to calibrate these parameter selections visually; after all, no numerical statistic is a perfect measure of what we actually see; at best, they are approximations of what we might potentially see. We have found it to be useful to have one experimenter calibrate the model parameters at a gross level, and then have another experimenter narrow in on the parameters which are




Figure 5: Density plot of the lack of fit statistic showing separation of selected difficulty levels: High (obvious curvature), Medium (noticeable curvature), and Low (almost linear). Each density plot is the result of 1000 simulations from a model $y_i = \alpha \cdot e^{\beta \cdot x_i + \epsilon_i} + \theta$, where $\epsilon_i \sim N(0, \sigma^2)$. α and θ were selected after manipulation of β and σ to ensure that all data generated had similar y ranges so as not to provide visual cues about model differences outside of the plot curvature.

visually reasonable within the selected range. Then, both examiners visually inspect a large number of plots generated using those parameters to get a sense for how difficult the task at hand is (this strategy is also described by Lu et al. [?]). At some point, all experimenters become so visually saturated with the nuances of the data generating mechanism that it may become necessary to “sanity check” the protocol with family members, friends, and colleagues. These informal surveys provide extremely useful feedback and can help to counteract the visual saturation of being immersed in the design of a visualization experiment for months at a time.

Visual Assessment is Critical

We cannot overstate the importance of visual assessment of your model stimuli, preferably with fresh eyes. We highly recommend performing a series of think-aloud pilot testing before deploying an experiment. In support of this assessment, we offer up a cautionary tale of our own experience: that of [?], where we designed an experiment to test which plot aesthetics promoted discovery of linear trends and/or clusters.

The experiment was a full 2x3x3 factorial exploration of three data generating parameters, with 3 replicates at each parameter combination (54 datasets) and 10 aesthetic combinations (for a total of 540 lineups). Each lineup had 20 different sub-panels, so we should have carefully visually inspected some 10800 different panels. As is evident from the fact that we’re telling this story as a cautionary tale, we missed a critical problem with our data-generating mechanism: when clusters were assigned to randomly generated data after the fact, we didn’t control the cluster size, leading to clusters of one or two points in relatively few sub-panels. This became particularly noticeable when bounding ellipses were added to the plot, as the method used to generate those ellipses required at least 3 points in the cluster. The missing boundary ellipse in the corresponding sub-panels escaped our notice during the stimuli proof-reading phase of the experiment, but did not escape the notice of our participants, who only needed to examine about 10 lineups each (around 200 panels). An example of one of the problematic lineups is shown in Figure 6: many participants selected panel 16 because of the missing ellipse; not a wrong choice, but certainly not the

effect we intended to test.




Figure 6: A lineup from Vanderplas & Hofmann [?]. Panel 10 shows the clustered target data and panel 17 shows the target data with a strong linear relationship; either of these target panels was the expected choice. Unfortunately, panel 16 has only two bounding ellipses shown, which is an unintentional difference that resulted from a faulty method for assigning clusters to null plots; many participants selected this panel instead of one of the target panels.

One reason why it is so difficult to generate sampling models for visual explorations is that our visual system is very, very good at finding differences and unexpected results. We reran the experiment using a different clustering method, and found that instead of noticing the number of ellipses, when that variable was removed participants instead cued in on the differences in size and shape of the ellipses formed using k-means clustering after the data generating procedure. That is, participants could still detect the artificial nature of the induced clusters using other features. While it can be difficult to get the data generating method right, it is essential to conducting visual experiments that generalize well beyond the effects shown in a single data set or phenomenon. The time and effort invested in this step at the outset of the experiment pays dividends when it allows for clear generalization of the experimental results to an entire statistical concept rather than a single data set.

Experimental Design Considerations

It would be difficult to develop a full data generating model without some idea of the experimental design: the basic structure of the parameters which are to be manipulated, how the users will be tested, and so on. These experimental design factors are fairly natural for scientists to accumulate over the course of imagining and planning an experiment. When conducting graphical tests, however, there are additional considerations beyond those taught in a standard experimental design course.

The primary human design factor to be aware of is that visual tasks and assessing statistical graphics can be extremely

cognitively taxing. In our experience, it is difficult to expect participants to evaluate more than about 15 statistical lineups in one sitting. It is easy to understand how this arises when we consider the concept of decision fatigue: each standard lineup has 20 panels, so there are 190 pairwise comparisons to evaluate between the plots of just a single lineup (though it is unlikely that participants actually sequentially evaluate each comparison in a combinatorial fashion). Tasks which are more interactive, such as ‘You Draw It’, may be somewhat easier for participants, but it is unlikely that participants would be willing to complete more than about 20 tasks in one sitting even with tasks that require fewer decisions and more engagement. At some point, participants’ interest in accuracy will decline, and the researcher is better off fitting the task within the window where participants are attentive and engaged.

It can be extremely helpful to include “practice” demonstrations of the task to show the basic process, logic, and reasoning. While it is tempting to make these tasks fully representative of the type of judgement which will be required of participants, practice tasks which are too close to the experimental task may bias participants; we have found that it works well to have a relatively easy practice task which utilizes a slightly different type of plot and/or type of data than what will be tested in the experiment. In cases which require interactivity, gif animations of the task being carried out are useful, as are additional visual cues, such as the yellow box used in the ‘You Draw It’ task¹ to indicate that there were points which were not completed.

While studies have found some relationship between lineup performance and demographic factors [?, ?], these differences are relatively small and unlikely to be important. Studies have also not found a strong relationship between visual task performance and participant recruitment method [?], perhaps in part because demographics which use social media sites and demographics completing online research tasks for pay overlap heavily. What is clear, however, is that some participants are very good at evaluating graphics and some participants are not; as a result, it can be extremely helpful to use random effects models with an effect for individual participants.

When we conduct visual graphics studies, we typically ask for basic demographic information from participants in order to assess how representative our sample is of the wider population; we typically do not analyze results by demographic variables in part because past studies have found so few differences in performance. It may be useful to ask a few more demographic questions about STEM education level for studies which ask more of participants from a mathematical standpoint; while lineup studies have not found strong associations with those variables, lineup studies also do not require participants to engage with the data presented in a chart in a way that requires higher-order mathematical reasoning. This has allowed us to make the argument to IRB that our research is exempt, as we do not collect enough demographic information to identify participants, however, this has occasionally come at some cost. Our recent study examining the use of log scales and exponential data was conducted using Prolific, which recruits participants from around the world; we required only that participants were fluent in English to participate. It was only after the experiment was completed that we realized that different countries introduce logarithms as a concept at different points during

primary and secondary education; it might be that individuals in some countries have much more experience with log scales than those educated in the United States. We mention this only to point out that while every experiment contains a few missed opportunities, it is worth giving careful thought to the demographic questions asked of participants and what information may be helpful during the analysis stage.

In statistical design terms, most of our studies involve some type of balanced incomplete block design, where participants are assigned to a subset of experimental conditions which allow for estimation of the full range of effects specified in the model. The particular structure of these designs depends heavily on the factorial structure of the study, but we typically arrange participants’ trials to ensure that participants see the same data set only once (where possible) and see as many different experimental conditions as possible. It is also important to reduce the impact of order effects using either e.g. Latin square designs or randomization where possible, but we recognize that this is not always feasible due to the need to maintain participant naivete during some portions of the experiment.

We have conducted visualization experiments using a wide variety of tools: custom web servers running experiments using PHP black magic, Qualtrics surveys for static graphics, and writing full Shiny applications that control every part of the experimental process from instructions to providing completion codes for payment. In our research, Shiny has provided the right balance between control over the experimental setting, procedure, etc. and the intricate details of web server management, however, this balance is likely different for every group and potentially for every experiment. Likewise, we have recruited participants using Amazon Mechanical Turk, Prolific, Reddit and other social media sites, and email; each recruitment method has trade-offs between cost, convenience, and control over demographic variables.

Pilot Testing and Quality Assurance

Once you have your experiment designed and your charts to test, the next step is an extensive round of pilot testing. The goal with pilot testing is to ensure that the experiment is set up precisely the way you want it, to ensure that your participants know what they are supposed to be doing and that any vague instructions are fixed, to provide time estimates for how long it will take to complete the task, and to provide you with some early data to examine for red flags. Our studies usually go through 2-3 rounds of pilot testing, with at least one of those rounds involving any relatives we can find that are less computer literate and will invariably figure out how to crash the testing application, and another round involving any and all coworkers and friends who may be available for 10-15 minutes during the work day. Pilot study samples do not need to be representative of any particular population, though we would caution against using exclusively visualization researchers or statisticians in your pilot sample because some issues are less likely to show up with knowledgeable participants. In some studies, each participant may see a new set of data, depending on how the study is designed. In these cases, we highly recommend saving all generated data to a database as well, so that it is possible to go back and examine exactly what happened and/or how things went wrong. Hard drive space is historically cheap relative to almost any other cost in an experiment; saving all of the data is a sensible measure.

One highly useful (but not strictly essential) component of an experiment that can be set up during the pilot testing

¹See a gif of testing with ‘You Draw It’ [here](#)

stage is a basic analysis script which summarizes all data collected to date visually. We have used such scripts in the past to produce automatically updating dashboards or web pages, allowing for real time or near-real time monitoring of data collection efforts. This provides an easy way to summarize completion of the experiment so that individuals can receive credit (if using services like MTurk or Prolific), but also allows interested participants to see the data if recruiting participants who are more likely to be interested in the results, as sometimes happens on social media sites. As data collection online can also happen extremely quickly (300 participants in <2h in our most recent Prolific experiment), this can provide an illusion of control over the deluge of data, allowing any issues to be spotted and resolved relatively quickly. If server load is a potential issue, it may also help to release batches of trials over a longer period of time in order to minimize the chance of having to make participants wait for others to complete the task before the server can handle additional connections².

Analyzing the Data

One constant with data analysis of these types of experiments is that no matter how carefully the experiment is planned, designed, and executed, there will be surprises. This is the dual curse and blessing of studying human perception: the visual system never quite works the way that we expect that it will, which provides endless fodder for science and occasionally complicates the data analysis.

Generalized Linear Mixed Models The see-value approach [?, ?, ?] is extremely useful for single lineups, but when a series of lineups that are part of a designed experiment are used, generalized linear mixed models are a much simpler way to summarize the overall effect of various manipulated factors [?, ?, ?]. This approach also works extremely well for psychometric experiments [?], as psychometric models can be easily fit into the framework of a generalized linear mixed model. In addition, similar model structures can be used to model accuracy, response time, and confidence, if all three types of information are collected from participants during the experiment.

Numerical Estimation There are additional considerations that should be expected when asking participants to estimate numerical quantities. Anchoring and rounding cause participant responses to cluster in ways that can bias statistical estimators, requiring methods designed for these types of data [?, ?, ?]. An alternative approach is to analyze the data graphically using e.g. a combination of density and rug geoms for point estimates (as shown in Figure 7), or spaghetti plots for individual response curves, with annotations to show anchor points. The advantage of this approach is that it can accommodate extremely messy data without requiring the extensive data cleaning and elimination of nonsensical responses that might be necessary to fit a statistical model.

Direct Interactions If participants are making predictions and/or fitting visual statistics, we have had success analyzing these responses by comparing the responses to results from a statistical model to determine how visual statistics differ from the numerical quantities derived mathematically. For instance, we calculate the deviation between participant

²This is the one major drawback to our preferred solution of self-hosting a Shiny server to handle data collection: the free version of Shiny server is limited to about 15 connections at any given time.




Figure 7: Density of participant estimates for the year in which the population reaches 4000. Colors are associated to scale - linear (blue) and log (orange) - and vertical lines indicate the true value based on the underlying model equation (black solid) and closest point value based on the simulated data set (black dashed). A jittered rug plot along the x -axis shows where participant estimates were made. The plot shows anchoring occurred to the closest point as shown by an increase in density around the dashed line. Density peaks occurred at whole values indicating rounding errors.

responses and the linear regression in ‘You Draw It’ experiments, then fitted generalized additive mixed models to summarize the results across different experimental conditions to assess how user-drawn predictions deviated from the statistical estimates. In other direct interactions, it may be useful to compare participant selections or annotations to closest points on the chart to assess anchoring behavior; for discrete selections, methods discussed in numerical estimation may also be useful.

Qualitative Responses In many cases, it is helpful to combine participants’ qualitative reasoning with their quantitative responses to designed graphical experiments. This approach provides useful context as to what participants use to make their decisions, and can be useful when assessing why unexpected responses occurred. We have used word clouds to show overall themes in participant responses successfully in [?]; when paired with an appropriate linear model it became clear that participants were fixating on unequal cluster size as a visual cue. In cases where participants are provided with additional utilities such as calculators and scratchpads, it can be useful to select responses from individual participants which illustrate the different types of calculations performed (but analyzing this data quantitatively can be difficult).

Conclusion

Testing features of visualization graphically using online platforms provides an incredibly powerful and efficient way to establish empirical guidelines for statistical graphics and visualization. There are nearly endless ways to combine web graphics, user interactions, and data collection to get insight into perception and use of graphics in practical settings. We have been continually surprised at the richness of the data collected in these experiments and the ability to combine qualitative and quantitative assessment to support conclusions that are both nuanced and of practical use when deciding how to design and present data using visualizations.